



Stony Brook  
University

Computer Science

# PROGRESS IMBALANCE IN MULTI-PROCESS PERFORMANCE

Arghya Bhattacharya

Algorithms Lab, Computer Science, Stony Brook University

# ARCHITECTURE OF MODERN COMPUTERS



Stony Brook  
University

Computer Science

- Multi-core
- Multi-threaded
- Time-shared

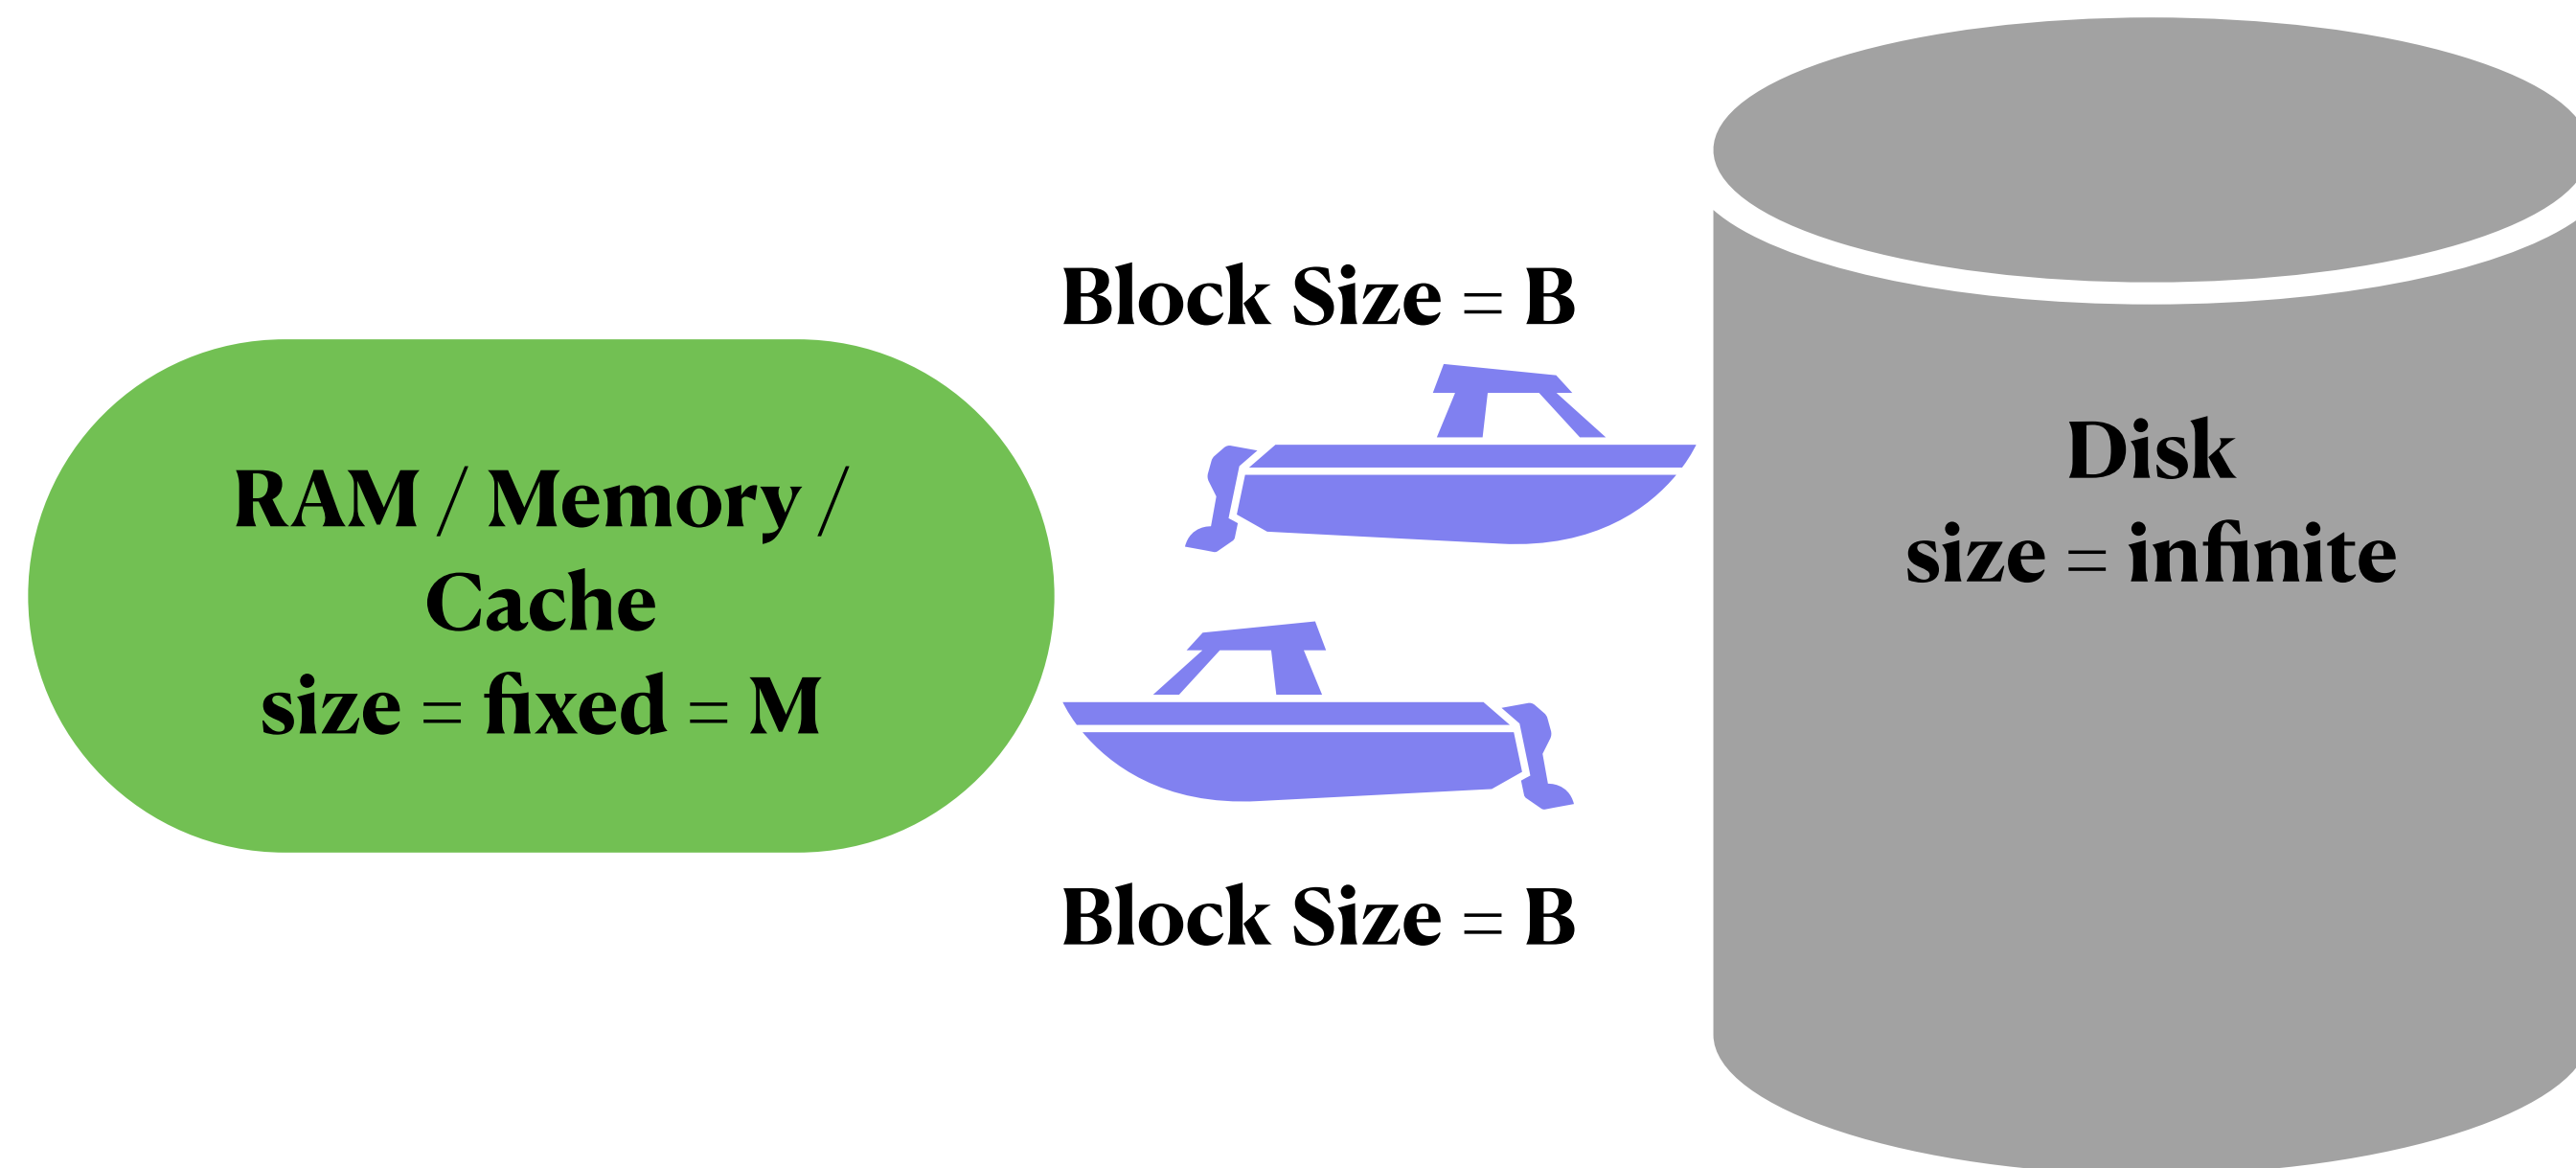
It is common that multiple programs share a memory

# DISK ACCESS MODEL



Stony Brook  
University

Computer Science



- Two-layered memory hierarchy
- Processing happens in the RAM
- Memory transfer (I/O) is reading from disk to RAM or vice versa
- This happens in chunks of size  $B$ .

[Aggarwal and Vitter, Communications of the ACM'88]

A program's running time is dependent on number of I/Os

# CACHE SIZE MATTERS



Stony Brook  
University

Computer Science

Example of I/O complexity:

Strassen's algorithm for Matrix Multiplication =  $\mathcal{O}\left(\frac{N^{\log_2 7}}{B\sqrt{M}}\right)$

External-memory Merge Sort =  $\mathcal{O}\left(\frac{N}{B} \log_{\frac{M}{B}} \frac{N}{B}\right)$

Cache available 

I/O made by a program 

# IMPORTANT TO UNDERSTAND THE CACHE BEHAVIOR

- DAM model assumes  $M$  is fixed.
- How is  $M$  shared among programs running concurrently?



Do the concurrent programs share a memory gracefully?



# DO THE CONCURRENT PROGRAMS SHARE A MEMORY GRACEFULLY?

**Experiment:** Running homogeneous instances concurrently

Homogeneous instances: copies of the same program with same problem size

**Observation:** A difference in the running time.

Progress imbalance: some program instances finish earlier than others.

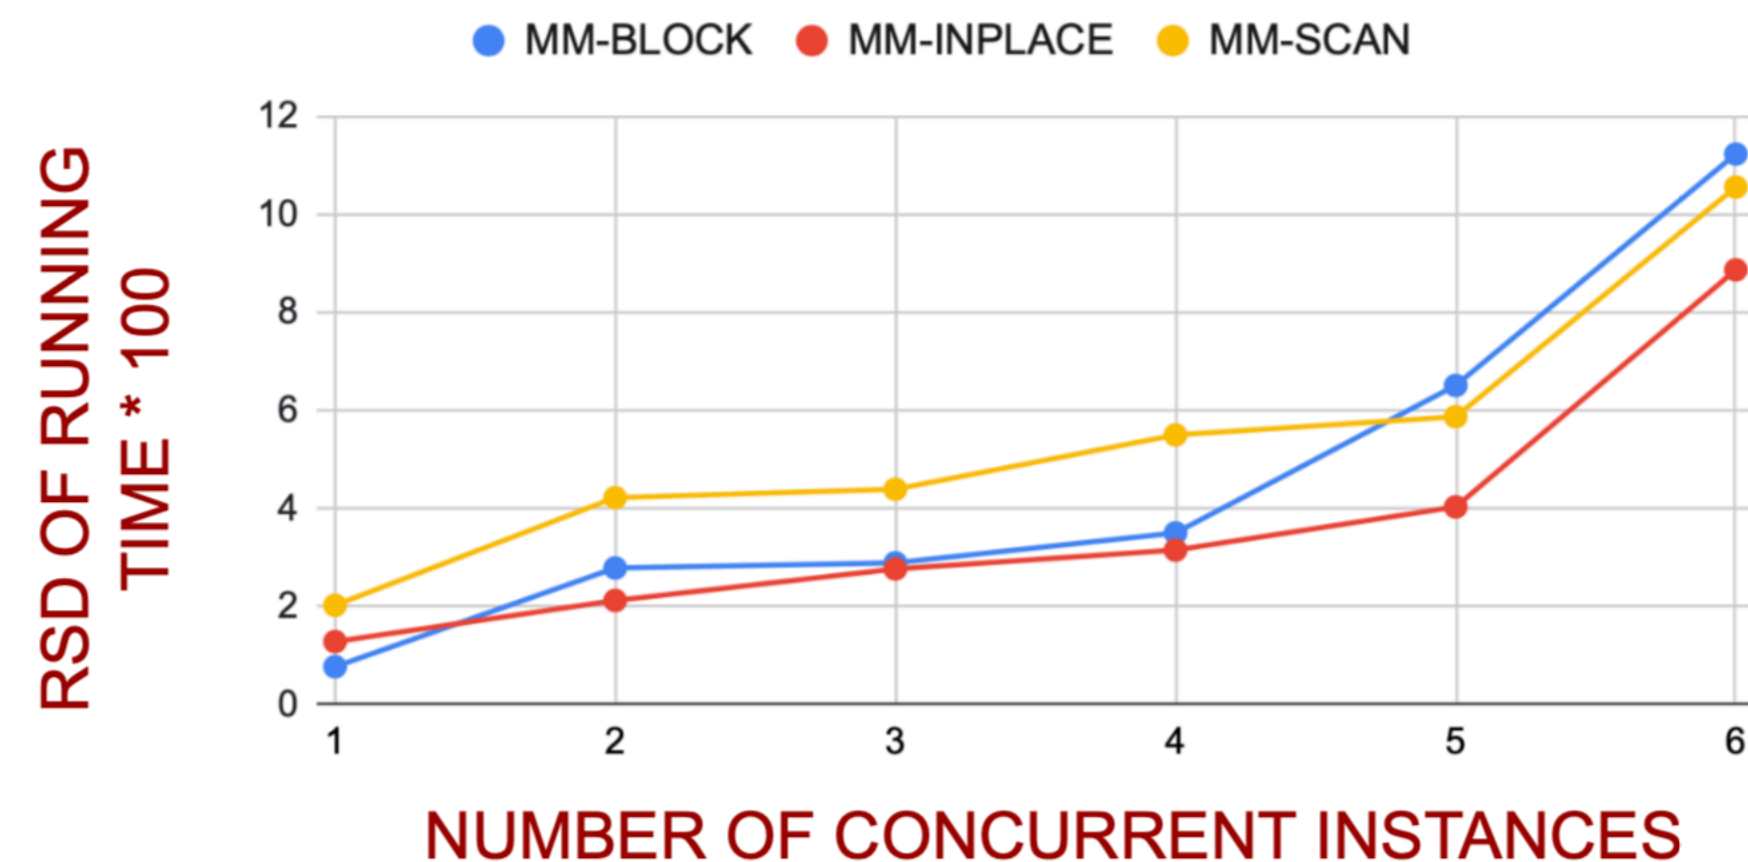


Number of concurrent processes ↑  
Progress imbalance ↑

# DO THE CONCURRENT PROGRAMS SHARE A MEMORY GRACEFULLY?

**Experiment:** Running a number of homogeneous instances concurrently

**Observation:** A progress imbalance of the instances



[Bender et al., SPAA'16]

Number of concurrent processes 

Progress imbalance 

# ETIOLOGY OF PROGRESS IMBALANCE

**Potential reason:** Imbalance in cache-sharing

**Experiment:** We run 6 concurrent instances.

Some instances have a transient stall (let's say  $x$  units of time) in the middle.

**What do we expect?**

- Instances will lose their share of cache when they stall.
- Again gain back their share when they resume.
- They finish with a delay of  $x$  units of time.

Is this what we observe?



# ETIOLOGY OF PROGRESS IMBALANCE

**Potential reason:** Imbalance in cache-sharing

**Experiment:** We run 6 concurrent instances.

Some instances have a transient stall (let's say  $x$  units of time) in the middle.

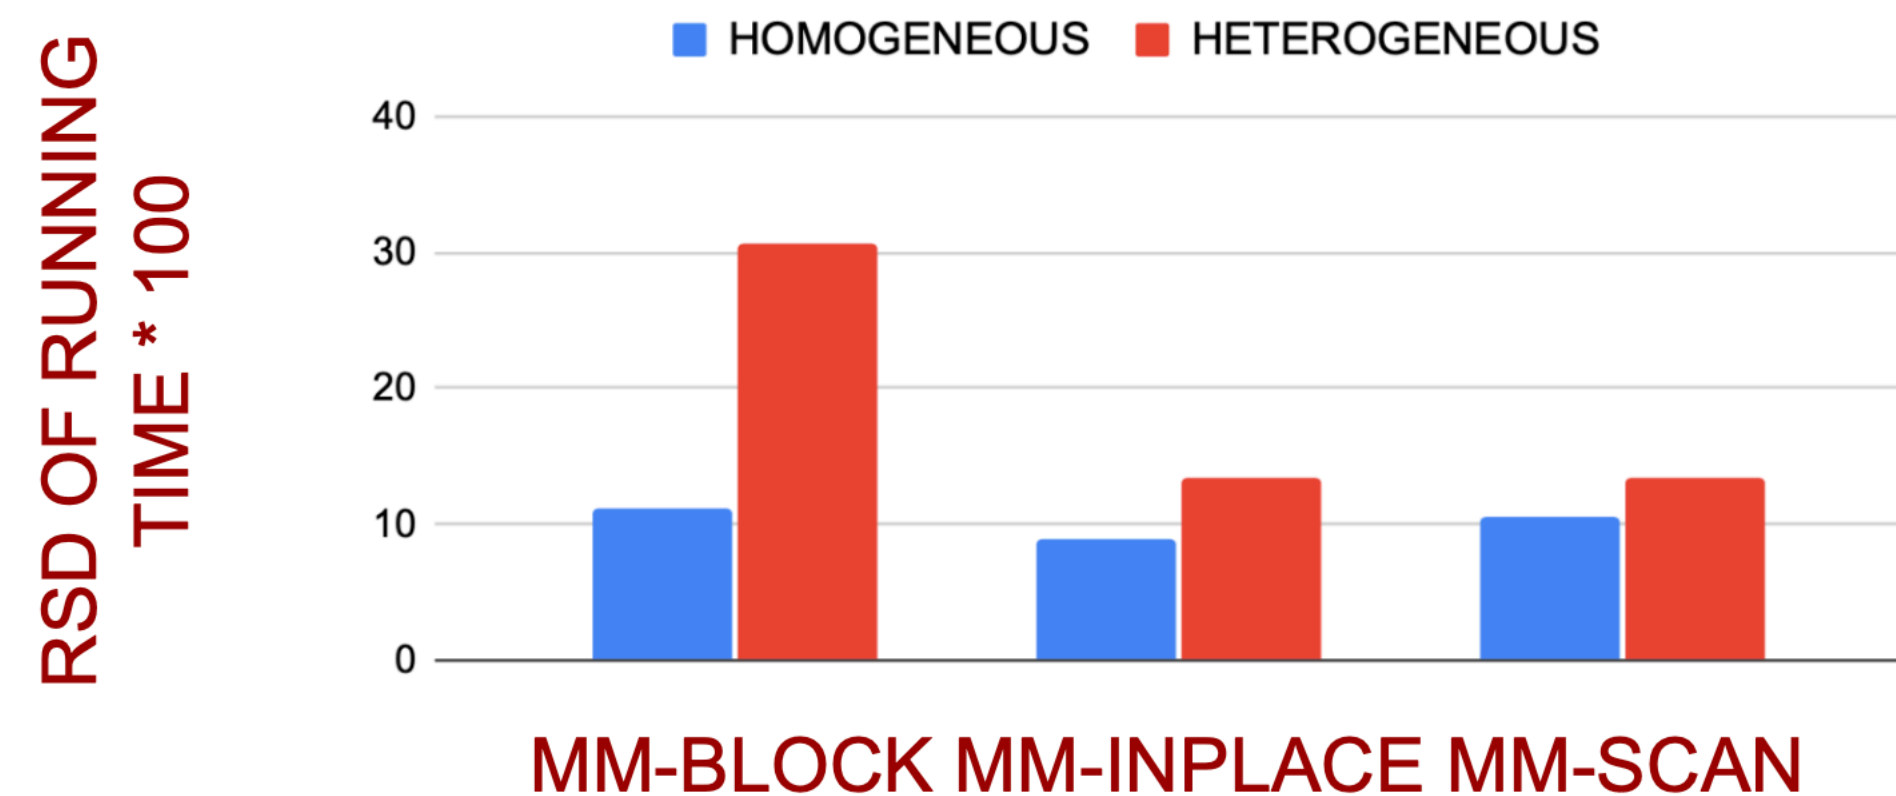
**What do we observe?**

- The instances with a transient stall lose their share of cache.
- They never gain that back.
- They finish much later than the other instances.
- This causes more progress imbalance.

Imbalance in cache-sharing leads to progress imbalance

# IMBALANCE IN CACHE-SHARING LEADS TO PROGRESS IMBALANCE

**Experiment:** We run 6 concurrent instances.  
Half of the instances have a transient stall (let's say  $x$  units of time) in the middle.



When multiple homogeneous programs run concurrently and share a cache, there's an imbalance in the cache-sharing, leading to a progress imbalance.

# FUTURE DIRECTIONS

- Quantify how badly different algorithms are affected by this phenomenon.
- Design algorithms that are susceptible.
- From the experiment, it seems if we randomly stall the programs, they may finish at the same time.



Stony Brook  
University

Computer Science

Thanks to:  
Prof. Michael A. Bender  
Prof. Rezaul A. Chowdhury

Thanks to all of you!